

K8s 03

1) Create a ClusterIP service for an Apache web server pod.

→ First create a pod with name apache-pod using below command:

```
kubectl run apache-pod --image=httpd --port=80
```

```
root@master:~# kubectl run apache-pod --image=httpd --port=80
pod/apache-pod created
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
apache-pod    1/1     Running   0           31s
```

→ now create a clusterIP service for apache-pod by using below command:

```
kubectl expose pod apache-pod --port=80 --target-port=80 --name=apache-
service (here type taken as default no need mention type in command)
```

```
root@master:~# kubectl expose pod apache-pod --port=80 --target-port=80 --name=apache-service
service/apache-service exposed
root@master:~# kubectl get services
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
apache-service ClusterIP    10.108.182.53 <none>        80/TCP    17s
```

→ publically not exposed its work only internally:

-- First do **kubectl get pods -o wide**

```
root@master:~# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE              NOMINATED NODE   READINESS GATES
apache-pod    1/1     Running   0           11m   192.168.123.137 ip-172-31-4-112   <none>           <none>
```

--then check it internally in master

```
curl 192.168.123.137
```

```
root@master:~# curl 192.168.123.137
<html><body><h1>It works!</h1></body></html>
root@master:~#
```

In worker-01

```
ubuntu@worker-01:~$ sudo -i
root@worker-01:~# curl 192.168.123.137
<html><body><h1>It works!</h1></body></html>
```

In worker-02

```
root@worker-02: ~
root@worker-02:~# curl 192.168.123.137
<html><body><h1>It works!</h1></body></html>
root@worker-02:~#
```

2) Expose an Nginx pod externally using a NodePort service.

→ First create nginx pod by using below command:

```
kubectl run nginx-pod --image=nginx --port=80
```

check it

```
root@master:~# kubectl run nginx-pod --image=nginx --port=80
pod/nginx-pod created
root@master:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
apache-pod	1/1	Running	0	22m
nginx-pod	1/1	Running	0	7s

→ now Expose an Nginx pod externally using a NodePort service by using below command:

```
kubectl expose pod nginx-pod --type=NodePort --name=nginx-service --port=80 -
-target-port=80
```

check it...Node port service has been created

```
root@master:~# kubectl expose pod nginx-pod --type=NodePort --name=nginx-service --port=80 --target-port=80
service/nginx-service exposed
root@master:~# kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
apache-service	ClusterIP	10.108.182.53	<none>	80/TCP	23m
firstpod-service	ClusterIP	10.111.36.119	<none>	80/TCP	19h
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	23h
nginx-service	NodePort	10.109.171.55	<none>	80:32679/TCP	25s


```
kubectl get services -o wide
```

```
root@master:~# kubectl get services -o wide
```


NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
apache-service	ClusterIP	10.108.182.53	<none>	80/TCP	26m	run=apache-pod
firstpod-service	ClusterIP	10.111.36.119	<none>	80/TCP	19h	app=myapp
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	23h	<none>
nginx-service	NodePort	10.109.171.55	<none>	80:32679/TCP	2m49s	run=nginx-pod

→ now check in browser:


Checked with master pubip:32679



Checked with worker-01 pubip:32679



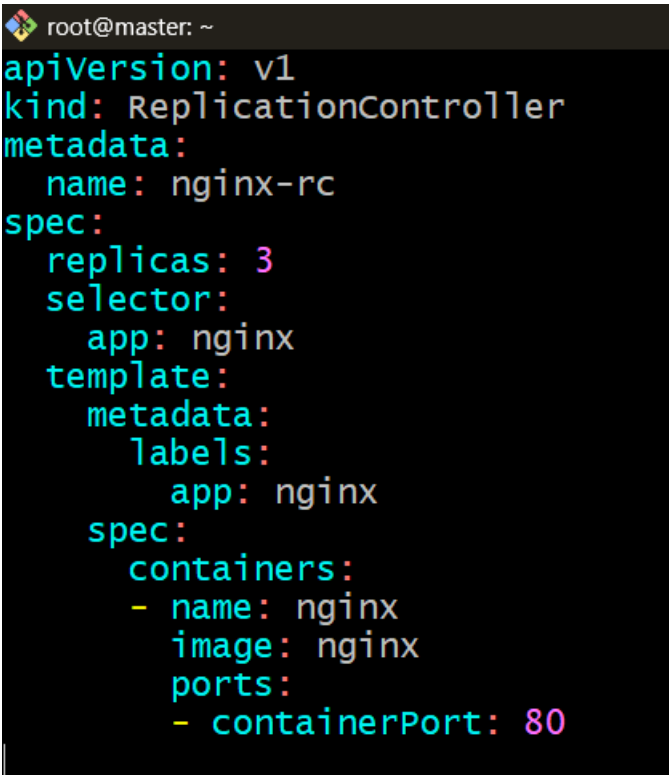
Checked with worker-02 pubip:32679



3) Deploy a ReplicationController to maintain 3 replicas of an Nginx pod.

→ Create ReplicationController using YAML file:

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx-rc
spec:
  replicas: 3
  selector:
    app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        ports:
        - containerPort: 80
```



A terminal window with a dark background and a light-colored cursor. The prompt is 'root@master: ~'. The terminal displays the following YAML configuration for a ReplicationController:

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx-rc
spec:
  replicas: 3
  selector:
    app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        ports:
        - containerPort: 80
```

→run the yaml file using below command to create rc:

```
kubectl apply -f nginx-rc.yaml
```

```
root@master:~# kubectl apply -f nginx-rc.yaml
replicationcontroller/nginx-rc created
```

--Rc created

```
kubectl get rc
```

```
root@master:~# kubectl get rc
NAME          DESIRED   CURRENT   READY   AGE
nginx-rc      3         3         3       18s
```

→now ReplicationController created to maintain 3 replicas of an Nginx pod:

Check with this command

```
kubectl get pods
```

```
root@master:~# kubectl get rc
NAME          DESIRED   CURRENT   READY   AGE
nginx-rc      3         3         3       18s
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx-rc-7pc2w 1/1     Running   0          94s
nginx-rc-c8d54 1/1     Running   0          94s
nginx-rc-jp8rr 1/1     Running   0          94s
```

→now delete one pod and check pods again created pods automatically:

```
root@master:~# kubectl delete pod nginx-rc-jp8rr
pod "nginx-rc-jp8rr" deleted
```

```
root@master:~# kubectl delete pod nginx-rc-jp8rr
pod "nginx-rc-jp8rr" deleted
root@master:~# ^C
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx-rc-7pc2w 1/1     Running   0          22m
nginx-rc-c8d54 1/1     Running   0          22m
nginx-rc-pvcs9 1/1     Running   0          9s
```

4) Scale the ReplicationController from 3 replicas to 5 replicas.

→ To scale a ReplicationController from 3 to 5 replicas, we can do it in two ways: using `kubectl` directly or by modifying the YAML file:

I choose `kubectl` command:

```
kubectl scale rc nginx-rc --replicas=5
```

```
root@master:~# kubectl scale rc nginx-rc --replicas=5
replicationcontroller/nginx-rc scaled
```

```
root@master:~# kubectl get rc
NAME          DESIRED  CURRENT  READY  AGE
nginx-rc      5        5        5      32m
```

→ check pods..2 more created:

```
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx-rc-7pc2w 1/1     Running   0          27m
nginx-rc-c8d54 1/1     Running   0          27m
nginx-rc-dgzss 1/1     Running   0          8s
nginx-rc-jdz8t 1/1     Running   0          8s
nginx-rc-pvcs9 1/1     Running   0          5m31s
```

5) Create a ReplicaSet to manage pods based on multiple labels (prod and test).

→ a ReplicaSet to manage pods based on multiple labels (prod and test):

apiVersion: apps/v1

kind: ReplicaSet

metadata:

name: my-replicaset

spec:

replicas: 3

selector:

matchLabels:

environment: prod

template:

metadata:

labels:

```
environment: prod
type: test
spec:
containers:
- name: my-app
image: my-app-image:latest
ports:
- containerPort: 80
```

```
root@master: ~
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      environment: prod
  template:
    metadata:
      labels:
        environment: prod
        type: test
    spec:
      containers:
      - name: my-app
        image: my-app-image:latest
        ports:
        - containerPort: 80
```

→run the yaml file:

```
kubectl create -f rs.yaml
```

created rs

```
root@master:~# kubectl create -f rs.yaml
replicaset.apps/my-replicaset created
```

→check labels:

```
kubectl get pods --show-labels
```

```
root@master:~# kubectl get pods --show-labels
NAME                                READY   STATUS    RESTARTS   AGE   LABELS
my-replicaset-521q7                 0/1    ImagePullBackOff    0      6s    environment=prod,type=test
my-replicaset-61x7f                 0/1    ImagePullBackOff    0      6s    environment=prod,type=test
my-replicaset-ckc5v                 0/1    ImagePullBackOff    0      6s    environment=prod,type=test
```

6) Deploy a ReplicaSet that excludes pods with the label backend.

→ **ReplicaSet that excludes pods with the label backend:**

apiVersion: apps/v1

kind: ReplicaSet

metadata:

name: firstrc

labels:

appname: testapp

spec:

replicas: 3

selector:

matchExpressions:

- key: env

operator: In

values:

- prod

- test

- key: type

operator: NotIn #ignore the pod with label as backend

values:

- backend

template:

metadata:

name: firstpod

labels:

env: prod

spec:

containers:

- name: firstcontainer

image: nginx

env:

- name: myname

```

root@master: ~
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: firstrc
  labels:
    appname: testapp
spec:
  replicas: 3
  selector:
    matchExpressions:
      - key: env
        operator: In
        values:
          - prod
          - test
      - key: type
        operator: NotIn    #ignore the pod with label as backend
        values:
          - backend
  template:
    metadata:
      name: firstpod
      labels:
        env: prod
    spec:
      containers:
        - name: firstcontainer
          image: nginx
          env:
            - name: myname

```

→run yaml file:

kubectl create -f rs.yaml

```

root@master: ~# kubectl create -f rs.yaml
replicaset.apps/my-replicaset created

```

→check labels:

--here It ignores backend type

```

root@master:~# kubectl get pods --show-labels

```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
firstrc-k865k	1/1	Running	0	6s	env=prod
firstrc-ncj6v	1/1	Running	0	6s	env=prod
firstrc-s2vgh	1/1	Running	0	6s	env=prod
my-replicaset-52lq7	0/1	ImagePullBackoff	0	9m21s	environment=prod,type=test
my-replicaset-6lx7f	0/1	ImagePullBackoff	0	9m21s	environment=prod,type=test
my-replicaset-ckc5v	0/1	ImagePullBackoff	0	9m21s	environment=prod,type=test

```

root@master:~# vi rs.yaml

```


7) Test load balancing across multiple pods using a NodePort service.

→yaml file to create multiple pods and create a nodeport service:

```
apiVersion: v1
kind: Service
metadata:
  name: echo-service
spec:
  type: NodePort
  selector:
    app: echo-server
  ports:
    - port: 80
      targetPort: 5678
      nodePort: 30080
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: echo-server
spec:
  replicas: 3
  selector:
    matchLabels:
      app: echo-server
  template:
    metadata:
      labels:
        app: echo-server
```

spec:

containers:

- name: echo

image: hashicorp/http-echo

args:

- "-text=Hello from pod: \${POD_NAME}"

env:

- name: POD_NAME

valueFrom:

fieldRef:

fieldPath: metadata.name

ports:

- containerPort: 5678

```

root@master: ~
apiVersion: v1
kind: Service
metadata:
  name: echo-service
spec:
  type: NodePort
  selector:
    app: echo-server
  ports:
    - port: 80
      targetPort: 5678
      nodePort: 30080
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: echo-server
spec:
  replicas: 3
  selector:
    matchLabels:
      app: echo-server
  template:
    metadata:
      labels:
        app: echo-server
    spec:
      containers:
        - name: echo
          image: hashicorp/http-echo
          args:
            - "-text=Hello from pod: $(POD_NAME)"
          env:
            - name: POD_NAME
              valueFrom:
                fieldRef:
                  fieldPath: metadata.name
          ports:
            - containerPort: 5678

```

→run the yaml:

kubectl apply -f nodeport-lb.yaml

```

root@master:~# vi nodeport-lb.yaml
root@master:~# kubectl apply -f nodeport-lb.yaml
service/echo-service created

```

→created service:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
apache-service	ClusterIP	10.108.182.53	<none>	80/TCP	27h
echo-service	NodePort	10.107.50.179	<none>	80:30080/TCP	24h


→ apply this command to get a node ip:

kubectl get nodes -o wide

```
root@master:~# kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-13-158                    Ready     <none>    26h   v1.29.15  172.31.13.158 <none>        Ubuntu 24.04.2 LTS   6.8.0-1026-aws   containerd://1.7
ip-172-31-4-112                    Ready     <none>    26h   v1.29.15  172.31.4.112  <none>        Ubuntu 24.04.2 LTS   6.8.0-1026-aws   containerd://1.7
master                              Ready     control-plane 26h   v1.29.15  172.31.1.127 <none>        Ubuntu 24.04.2 LTS   6.8.0-1026-aws   containerd://1.7
```

→ now access it on browser worker-01 pub ip:nodeport:

We can see load balancing in action....



8) Delete a ReplicationController without affecting the running pods

→ first create a replica set:

```
root@master:~# vi nginx-rc.yaml
root@master:~# kubectl apply -f nginx-rc.yaml
replicationcontroller/nginx-rc created
```

→ check rc:

```
root@master:~# kubectl get rc
NAME          DESIRED   CURRENT   READY   AGE
nginx-rc      3         3         3       94s
```

→ check pods:

```
root@master:~# kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
nginx-rc-6t9qs      1/1     Running   0           11s
nginx-rc-8h7jn      1/1     Running   0           11s
nginx-rc-q5fgd      1/1     Running   0           11s
```

→now Delete a ReplicationController without affecting the running pods using below command:

```
kubectl delete rc --cascade=false nginx-rc
```

now check Rc

```
root@master:~# kubectl get rc
No resources found in default namespace.
```

Deleted rc

Now check pods

```
root@master:~# kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
nginx-rc-6t9qs       1/1     Running   0           2m16s
nginx-rc-8h7jn       1/1     Running   0           2m16s
nginx-rc-q5fgd       1/1     Running   0           2m16s
```

Pods are not deleted.....