

K8s 04

1) Deploy an application using a Deployment with 3 replicas and a rolling update strategy.

→ create a yaml file:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: firstdeployment
  labels:
    appname: testapp
```

```
spec:
  replicas: 3
  selector:
    matchExpressions:
      - key: env
        operator: In
        values:
          - prod
  template:
    metadata:
      name: firstpod
      labels:
        env: prod
    spec:
      containers:
        - name: firstcontainer
          image: nginx
          env:
            - name: myname
```

```

ubuntu@master: ~
apiVersion: apps/v1
kind: Deployment
metadata:
  name: firstdeployment
  labels:
    appname: testapp
spec:
  replicas: 3
  selector:
    matchExpressions:
      - key: env
        operator: In
        values:
          - prod
  template:
    metadata:
      name: firstpod
      labels:
        env: prod
    spec:
      containers:
        - name: firstcontainer
          image: nginx
          env:
            - name: myname

```

→run yaml file:

```

ubuntu@master:~$ kubectl apply -f deployment.yaml
deployment.apps/firstdeployment created

```

→created Deployment with 3 replicas(pods):

```

ubuntu@master:~$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
firstdeployment     3/3     3            3           21s
ubuntu@master:~$ kubectl rollout status deployment/firstdeployment.yaml
Error from server (NotFound): deployments.apps "firstdeployment.yaml" not found
ubuntu@master:~$ kubectl rollout status deployment/firstdeployment
deployment "firstdeployment" successfully rolled out
ubuntu@master:~$ kubectl rollout history deployment/firstdeployment
deployment.apps/firstdeployment
REVISION  CHANGE-CAUSE
1          <none>

```

--creted 3 pods

```

root@master:~# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
firstdeployment-c6b5c8445-2z5pr     1/1     Running   0          11m
firstdeployment-c6b5c8445-6hv9v     1/1     Running   0          11m
firstdeployment-c6b5c8445-rjtmw     1/1     Running   0          11m

```

→ Deployed an application using a Deployment with 3 replicas and a rolling update strategy:

```
ubuntu@master:~$ kubectl describe deployments --recursive
Name: firstdeployment
Namespace: default
CreationTimestamp: Wed, 16 Apr 2025 09:03:03 +0000
Labels: appname=testapp
Annotations: deployment.kubernetes.io/revision: 1
Selector: env in (prod)
Replicas: 3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels: env=prod
  Containers:
    firstcontainer:
      Image: nginx
```

2) Configure a Deployment with a Recreate strategy and observe the downtime.

→ create yaml file:

```
root@master: ~  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: firstdeployment  
  labels:  
    appname: testapp  
spec:  
  replicas: 3  
  strategy:  
    type: Recreate  
  selector:  
    matchExpressions:  
      - key: env  
        operator: In  
        values:  
          - prod  
  template:  
    metadata:  
      name: firstpod  
      labels:  
        env: prod  
    spec:  
      containers:  
        - name: firstcontainer  
          image: nginx  
          env:  
            - name: myname
```

→ run the yaml file:

```
root@master:~# kubectl get deployments  
NAME                READY   UP-TO-DATE   AVAILABLE   AGE  
firstdeployment     3/3     3            3           15m
```

→ Configured a Deployment with a Recreate strategy:

```
root@master:~# kubectl describe deployments --recursive
Name:          firstdeployment
Namespace:     default
CreationTimestamp: Wed, 16 Apr 2025 09:26:35 +0000
Labels:        appname=testapp
Annotations:    deployment.kubernetes.io/revision: 1
Selector:      env in (prod)
Replicas:      3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:  Recreate
MinReadySeconds: 0
Pod Template:
  Labels:  env=prod
  Containers:
    firstcontainer:
      Image:  nginx
```

→ observed the downtime:

```
root@master:~# kubectl get pods -w
NAME                                READY    STATUS    RESTARTS   AGE
firstdeployment-c6b5c8445-2z5pr      1/1      Running   0           20m
firstdeployment-c6b5c8445-6hv9v      1/1      Running   0           20m
firstdeployment-c6b5c8445-rjtmw      1/1      Running   0           20m
firstdeployment-c6b5c8445-2z5pr      1/1      Terminating   0           23m
firstdeployment-c6b5c8445-rjtmw      1/1      Terminating   0           23m
firstdeployment-c6b5c8445-6hv9v      1/1      Terminating   0           23m
firstdeployment-c6b5c8445-2z5pr      1/1      Terminating   0           23m
firstdeployment-c6b5c8445-rjtmw      1/1      Terminating   0           23m
firstdeployment-c6b5c8445-6hv9v      1/1      Terminating   0           23m
firstdeployment-c6b5c8445-2z5pr      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-rjtmw      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-6hv9v      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-2z5pr      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-rjtmw      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-6hv9v      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-2z5pr      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-rjtmw      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-6hv9v      0/1      Terminating   0           23m
firstdeployment-c6b5c8445-2z5pr      0/1      Pending        0           0s
firstdeployment-c6b5c8445-rjtmw      0/1      Pending        0           0s
firstdeployment-c6b5c8445-6hv9v      0/1      Pending        0           0s
firstdeployment-c6b5c8445-2z5pr      0/1      Pending        0           0s
firstdeployment-c6b5c8445-rjtmw      0/1      Pending        0           0s
firstdeployment-c6b5c8445-6hv9v      0/1      Pending        0           0s
firstdeployment-c6b5c8445-2z5pr      0/1      ContainerCreating   0           0s
firstdeployment-c6b5c8445-rjtmw      0/1      ContainerCreating   0           0s
firstdeployment-c6b5c8445-6hv9v      0/1      ContainerCreating   0           0s
firstdeployment-c6b5c8445-2z5pr      0/1      ContainerCreating   0           1s
firstdeployment-c6b5c8445-rjtmw      0/1      ContainerCreating   0           1s
firstdeployment-c6b5c8445-6hv9v      0/1      ContainerCreating   0           1s
firstdeployment-c6b5c8445-2z5pr      1/1      Running           0           1s
firstdeployment-c6b5c8445-rjtmw      1/1      Running           0           1s
firstdeployment-c6b5c8445-6hv9v      1/1      Running           0           1s
firstdeployment-c6b5c8445-2z5pr      1/1      Running           0           2s
```

```
ubuntu@master:~$ sudo -i
root@master:~# kubectl apply -f deployment.yaml
deployment.apps/firstdeployment configured
root@master:~# kubectl delete deployments
error: resource(s) were provided, but no name was specified
root@master:~# kubectl delete firstdeployment
error: the server doesn't have a resource type "firstdeployment"
root@master:~# kubectl delete deployment firstdeployment
deployment.apps "firstdeployment" deleted
root@master:~# kubectl apply -f deployment.yaml
deployment.apps/firstdeployment created
root@master:~#
```

```
^Croot@master:~# kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
firstdeployment-c6b5c8445-97zwz      1/1      Running   0           3m2s
firstdeployment-c6b5c8445-t5w7d      1/1      Running   0           3m2s
firstdeployment-c6b5c8445-wr6fz      1/1      Running   0           3m2s
```

3) Update an existing Deployment and perform a rollback to the previous version.

→first check revisions:

```
root@master:~# kubectl rollout history deployment firstdeployment
deployment.apps/firstdeployment
REVISION  CHANGE-CAUSE
1          <none>
2          nginx changed to httpd
```

→use this command to rollback to the previous version:

kubectl rollout undo --to-revision=2 deployment <deployment_name>

```
root@master:~# kubectl rollout undo --to-revision=1 deployment firstdeployment
deployment.apps/firstdeployment rolled back
```

→now check history:

```
root@master:~# kubectl rollout history deployment firstdeployment
deployment.apps/firstdeployment
REVISION  CHANGE-CAUSE
2          nginx changed to httpd
3          <none>
```

4) Modify a Deployment to add resource requests and limits for CPU and memory.

→ Modified a Deployment to add resource requests and limits for CPU and memory:

```
root@master: ~  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: firstdeployment  
  labels:  
    appname: testapp  
  
spec:  
  replicas: 6  
  strategy:  
    type: Recreate  
  
  selector:  
    matchExpressions:  
      - key: env  
        operator: In  
        values:  
          - prod  
  
  template:  
    metadata:  
      name: firstpod  
      labels:  
        env: prod  
  
    spec:  
      containers:  
        - name: firstcontainer  
          image: nginx  
          env:  
            - name: myname  
              value: "your_name_here"  
          resources:  
            requests:  
              memory: "128Mi"  
              cpu: "250m"  
            limits:  
              memory: "256Mi"  
              cpu: "500m"
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:
```

name: firstdeployment

labels:

appname: testapp

spec:

replicas: 6

strategy:

type: Recreate

selector:

matchExpressions:

- key: env

operator: In

values:

- prod

template:

metadata:

name: firstpod

labels:

env: prod

spec:

containers:

- name: firstcontainer

image: nginx

env:

- name: myname

value: "your_name_here"

resources:

requests:

memory: "128Mi"

cpu: "250m"

limits:

memory: "256Mi"

cpu: "500m"

→run the yaml file:

```
root@master:~# kubectl apply -f deployment.yaml
deployment.apps/firstdeployment configured
```

→check with describe command:

```
kubectl describe pod firstdeployment-6798b47df7-75phj
```

Restart Count:	0
Limits:	
cpu:	500m
memory:	256Mi
Requests:	
cpu:	250m
memory:	128Mi

5) Create a Deployment with MaxSurge and MaxUnavailable configurations.

→yaml file to Create a Deployment with MaxSurge and MaxUnavailable configurations:

```
root@master: ~  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: firstdeployment  
  labels:  
    appname: testapp  
spec:  
  replicas: 6  
  minReadySeconds: 30  
  strategy:  
    rollingUpdate:  
      maxSurge: 2  
      maxUnavailable: 1  
  selector:  
    matchExpressions:  
      - key: env  
        operator: In  
        values:  
          - prod  
  template:  
    metadata:  
      name: firstpod  
      labels:  
        env: prod  
    spec:  
      containers:  
        - name: firstcontainer  
          image: nginx  
          env:  
            - name: myname
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: firstdeployment
  labels:
    appname: testapp
```

```
spec:
  replicas: 6
  minReadySeconds: 30
  strategy:
    rollingUpdate:
      maxSurge: 2
      maxUnavailable: 1
  selector:
    matchExpressions:
      - key: env
        operator: In
        values:
          - prod
  template:
    metadata:
      name: firstpod
      labels:
        env: prod
    spec:
      containers:
        - name: firstcontainer
          image: nginx
          env:
            - name: myname
```

→run the yaml file:

```
root@master:~# kubectl apply -f deployment.yaml
deployment.apps/firstdeployment configured
```

→check MaxSurge and MaxUnavailable configurations:

```
root@master:~# kubectl describe deployments
Name: firstdeployment
Namespace: default
CreationTimestamp: Wed, 16 Apr 2025 09:26:35 +0000
Labels: appname=testapp
Annotations: deployment.kubernetes.io/revision: 5
Selector: env in (prod)
Replicas: 6 desired | 6 updated | 8 total | 5 available | 3 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 30
RollingUpdateStrategy: 1 max unavailable, 2 max surge
```

--as MaxSurge 2 first created 8 pods then 2stoped finally 6 pods only created:

```
root@master:~# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
firstdeployment-6798b47df7-nxzr7    1/1     Running   0           17m
firstdeployment-6798b47df7-xg8rd    1/1     Running   0           17m
firstdeployment-c6b5c8445-6pzlr     1/1     Running   0           54s
firstdeployment-c6b5c8445-bmfpm     1/1     Running   0           22s
firstdeployment-c6b5c8445-c8chc     1/1     Running   0           54s
firstdeployment-c6b5c8445-cz2l6     1/1     Running   0           54s
firstdeployment-c6b5c8445-gnxx5     1/1     Running   0           22s
firstdeployment-c6b5c8445-gxw9t     1/1     Running   0           22s
root@master:~# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
firstdeployment-6798b47df7-nxzr7    0/1     Terminating   0           17m
firstdeployment-6798b47df7-xg8rd    0/1     Terminating   0           17m
firstdeployment-c6b5c8445-6pzlr     1/1     Running        0           65s
firstdeployment-c6b5c8445-bmfpm     1/1     Running        0           33s
firstdeployment-c6b5c8445-c8chc     1/1     Running        0           65s
firstdeployment-c6b5c8445-cz2l6     1/1     Running        0           65s
firstdeployment-c6b5c8445-gnxx5     1/1     Running        0           33s
firstdeployment-c6b5c8445-gxw9t     1/1     Running        0           33s
root@master:~# kubectl get deployments
NAME            READY   UP-TO-DATE   AVAILABLE   AGE
firstdeployment 6/6      6             6           62m
```

6) Set up a Deployment with a custom revision history limit.

→yaml file Set up a Deployment with a custom revision history limit:

apiVersion: apps/v1

kind: Deployment

metadata:

name: my-first-deployment

labels:

app: myapp

spec:

replicas: 5

revisionHistoryLimit: 15

selector:

matchLabels:

app: myapp

template:

metadata:

```
labels:
  app: myapp
spec:
  containers:
  - name: myapp-container
    image: nginx:latest
    ports:
    - containerPort: 80
  resources:
    requests:
      memory: "128Mi"
      cpu: "200m"
    limits:
      memory: "256Mi"
      cpu: "500m"
```

```
root@master: ~
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-first-deployment
  labels:
    app: myapp
spec:
  replicas: 5
  revisionHistoryLimit: 15
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
      - name: myapp-container
        image: nginx:latest
        ports:
        - containerPort: 80
      resources:
        requests:
          memory: "128Mi"
          cpu: "200m"
        limits:
          memory: "256Mi"
          cpu: "500m"
```

→run the yaml file:

```
root@master:~# kubectl apply -f deployment.yaml
deployment.apps/my-first-deployment created
```

→created deployment:

```
root@master:~# kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
firstdeployment      6/6     6            6           65m
my-first-deployment  5/5     5            5           13s
```

→check custom revision history limit with below command:

```
kubectl get deployment my-first-deployment -o yaml
```

```
uid: 85867e20-7d4a-4888-b268-177d3f1dc
spec:
  progressDeadlineSeconds: 600
  replicas: 5
  revisionHistoryLimit: 15
  selector:
    matchLabels:
      app: myapp
```

7) Pause a Deployment during an update, and then resume it.

→to Pause Deployment name “my-first-deployment” use below command:

```
kubectl rollout pause deployment my-first-deployment
```

```
root@master:~# kubectl rollout pause deployment my-first-deployment
deployment.apps/my-first-deployment paused
```

→to resume a deployment use below command:

```
kubectl rollout resume deployment my-first-deployment
```

```
root@master:~# kubectl rollout resume deployment my-first-deployment
deployment.apps/my-first-deployment resumed
```

8) Create a pod using resource requests for memory and CPU, and observe how the scheduler assigns it to a node.

→ create yaml file to Create a pod using resource requests for memory and CPU:

apiVersion: v1

kind: Pod

metadata:

name: resource-pod

spec:

containers:

- name: nginx

image: nginx

resources:

requests:

memory: "100Mi"

cpu: "200m"

limits:

memory: "200Mi"

cpu: "500m"

→run yaml file:

--created a pod

```
root@master:~# kubectl apply -f resource-pod.yaml
pod/resource-pod created
```

→check pods:

```
kubectl get pods -o wide
```

--Now we can see which node it was placed on by the scheduler

```
root@master:~# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE              NOMINATED NODE   READINESS GATES
resource-pod   1/1     Running   0           18s   192.168.201.249   ip-172-31-13-158   <none>           <none>
```

→Look under the Events section we'll see lines like:

```
Events:
  Type     Reason      Age   From              Message
  ----     -
  Normal   Scheduled   2m14s   default-scheduler   Successfully assigned default/resource-pod to ip-172-31-13-158
  Normal   Pulling     2m14s   kubelet            Pulling image "nginx"
  Normal   Pulled      2m14s   kubelet            Successfully pulled image "nginx" in 230ms (230ms including waiting for new image metadata)
  Normal   Created     2m14s   kubelet            Created container: nginx
  Normal   Started     2m14s   kubelet            Started container nginx
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
```

→Check Node Allocatable Resources with this command:

```
Kubectl describe node allocated node
```

```
root@master:~# kubectl describe node ip-172-31-13-158
Name:          ip-172-31-13-158
Roles:         <none>
Labels:        beta.kubernetes.io/arch=amd64
               beta.kubernetes.io/os=linux
               kubernetes.io/arch=amd64
               kubernetes.io/hostname=ip-172-31-13-158
               kubernetes.io/os=linux
Annotations:   kubeadm.alpha.kubernetes.io/cri-socket: unix:///var/run/contain
               node.alpha.kubernetes.io/ttl: 0
               projectcalico.org/IPv4Address: 172.31.13.158/20
               projectcalico.org/IPv4IPtunnelAddr: 192.168.201.192
               volumes.kubernetes.io/controller-managed-attach-detach: true
```

--based on below resource-limits and requests the scheduler assigns pod to a node...

Allocated resources:		
(Total limits may be over 100 percent, i.e., overcommitted.)		
Resource	Requests	Limits
-----	-----	-----
cpu	600m (30%)	500m (25%)
memory	134Mi (3%)	298Mi (7%)
ephemeral-storage	0 (0%)	0 (0%)
hugepages-2Mi	0 (0%)	0 (0%)