

## K8s 05

- 1) Create a namespace dev-environment and apply a resource-based quota that restricts the number of pods to 3 and services to 2.

→ Create the dev-environment Namespace by using below command:

```
kubectl create namespace dev-environment
```

```
ubuntu@master:~$ kubectl create namespace dev-environment
namespace/dev-environment created
```

--check ns..its created

```
ubuntu@master:~$ kubectl get ns
NAME                STATUS   AGE
default             Active   3d23h
dev-environment     Active   14s
kube-node-lease     Active   3d23h
kube-public         Active   3d23h
kube-system         Active   3d23h
tigera-operator     Active   3d22h
```

→YAML File to create a resource-based quota that restricts the number of pods to 3 and services to 2:

```
ubuntu@master: ~
apiVersion: v1
kind: ResourceQuota
metadata:
  name: dev-quota
  namespace: dev-environment
spec:
  hard:
    pods: "3"
    services: "2"
```

--run Yaml file

```
ubuntu@master:~$ vi resource-quota.yaml
ubuntu@master:~$ kubectl apply -f resource-quota.yaml
resourcequota/dev-quota created
```

--check

```
ubuntu@master:~$ kubectl get resourcequota -n dev-environment
NAME          AGE   REQUEST              LIMIT
dev-quota    20s   pods: 0/3, services: 0/2
```

→described a resource-based quota that restricts the number of pods to 3 and services to 2:

### kubectl describe ns dev-environment

```
ubuntu@master:~$ kubectl describe ns dev-environment
Name:          dev-environment
Labels:        kubernetes.io/metadata.name=dev-environment
Annotations:   <none>
Status:        Active

Resource Quotas
  Name:      dev-quota
  Resource   Used   Hard
  ---
  pods       0     3
  services   0     2

No LimitRange resource.
ubuntu@master:~$
```

→ now checked that when we try to create more than 3 pods got error so its restricted by resource-based quota:

```
ubuntu@master:~$ kubectl run mypod1 --image=nginx -n dev-environment
kubectl run mypod2 --image=nginx -n dev-environment
kubectl run mypod3 --image=nginx -n dev-environment
kubectl run mypod4 --image=nginx -n dev-environment
pod/mypod1 created
pod/mypod2 created
pod/mypod3 created
Error from server (Forbidden): pods "mypod4" is forbidden: exceeded quota: dev-quota, requested: pods=1, used: pods=3, limited: pods=3
```

## 2) Create a pod in the prod-environment namespace with 0.2 CPU and 200Mi memory requests, and 0.5 CPU and 500Mi memory limits.

→ first create namespace by below command:

### kubectl create ns prod-environment

```
root@master:~# kubectl create ns prod-environment
namespace/prod-environment created
```

--created:

```
root@master:~# kubectl get ns
NAME                STATUS      AGE
default             Active      3d23h
dev-environment     Terminating 24m
kube-node-lease     Active      3d23h
kube-public         Active      3d23h
kube-system         Active      3d23h
prod-environment    Active      6s
tigera-operator     Active      3d22h
```

→ yaml file to Create a pod in the prod-environment namespace with 0.2 CPU and 200Mi memory requests, and 0.5 CPU and 500Mi memory limits:

```
root@master: ~  
apiVersion: v1  
kind: Pod  
metadata:  
  name: my-pod  
  namespace: prod-environment  
spec:  
  containers:  
  - name: my-container  
    image: nginx  
    resources:  
      requests:  
        cpu: "0.2"  
        memory: "200Mi"  
      limits:  
        cpu: "0.5"  
        memory: "500Mi"
```

--run yaml

```
root@master:~# kubectl apply -f prod-pod.yaml  
pod/my-pod configured
```

--check

```
root@master:~# kubectl get pods -n prod-environment  
NAME       READY   STATUS    RESTARTS   AGE  
my-pod     1/1     Running   1 (106s ago)  31m
```

→describe pod:

```
root@master:~# kubectl describe pod my-pod -n prod-environment  
Name:         my-pod  
Namespace:    prod-environment  
Priority:      0
```

```
Ready:        True  
Restart Count: 1  
Limits:  
  cpu:        500m  
  memory:     500Mi  
Requests:  
  cpu:        200m  
  memory:     200Mi  
Environment:  <none>
```

- 3) In the staging-environment namespace, set a LimitRange that assigns default CPU and memory limits (300m CPU, 600Mi memory) and applies a minimum and maximum CPU.

→create name space:

```
root@master:~# kubectl create ns staging-environment
namespace/staging-environment created
root@master:~# kubectl get ns
```

| NAME                | STATUS      | AGE   |
|---------------------|-------------|-------|
| default             | Active      | 4d    |
| dev-environment     | Terminating | 71m   |
| kube-node-lease     | Active      | 4d    |
| kube-public         | Active      | 4d    |
| kube-system         | Active      | 4d    |
| prod-environment    | Terminating | 47m   |
| staging-environment | Active      | 6s    |
| tigera-operator     | Active      | 3d23h |

→yaml file to create limit range:

```
root@master: ~
apiVersion: v1
kind: LimitRange
metadata:
  name: default-limits
  namespace: staging-environment
spec:
  limits:
  - default:
      cpu: "300m"
      memory: "600Mi"
    defaultRequest:
      cpu: "300m"
      memory: "600Mi"
    min:
      cpu: "100m"
      memory: "200Mi"
    max:
      cpu: "1000m"
      memory: "2Gi"
    type: Container
```

--run yaml file

```
root@master:~# kubectl apply -f limitrange.yaml
limitrange/default-limits created
```

→describe:

```
root@master:~# kubectl describe ns staging-environment
Name:          staging-environment
Labels:        kubernetes.io/metadata.name=staging-environment
Annotations:   <none>
Status:        Active
```

No resource quota.

| Resource Limits |          |       |     |                 |               |                         |
|-----------------|----------|-------|-----|-----------------|---------------|-------------------------|
| Type            | Resource | Min   | Max | Default Request | Default Limit | Max Limit/Request Ratio |
| Container       | cpu      | 100m  | 1   | 300m            | 300m          | -                       |
| Container       | memory   | 200Mi | 2Gi | 600Mi           | 600Mi         | -                       |

4) Create a pod and a NodePort service in the default namespace, then create another pod in the test namespace and communicate between them using Service DNS.

→ yaml file to create a pod in default namespace:

```
root@master: ~
apiVersion: v1
kind: Pod
metadata:
  name: httpd-pod
  namespace: default
spec:
  containers:
  - name: httpd
    image: httpd
    ports:
    - containerPort: 80
```

--run the yaml then pod created in default namespace:

```
root@master:~# kubectl apply -f httpd-pod.yaml
pod/httpd-pod created
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
httpd-pod     1/1     Running   0           23s
testpod       1/1     Running   0           49m
```

→now create the nodeport service:

```
root@master: ~  
apiVersion: v1  
kind: Service  
metadata:  
  name: httpd-service  
  namespace: default  
spec:  
  selector:  
    app: httpd  
  ports:  
    - protocol: TCP  
      port: 80  
      targetPort: 80  
      nodePort: 30081  
  type: NodePort
```

→run yaml then created nodeport service:

```
root@master:~# vi httpd-service.yaml  
root@master:~# kubectl apply -f httpd-service.yaml  
service/httpd-service created
```

```
root@master:~# kubectl get services
```

| NAME             | TYPE      | CLUSTER-IP    | EXTERNAL-IP | PORT(S)      | AGE   |
|------------------|-----------|---------------|-------------|--------------|-------|
| apache-service   | ClusterIP | 10.108.182.53 | <none>      | 80/TCP       | 3d1h  |
| echo-service     | NodePort  | 10.107.50.179 | <none>      | 80:30080/TCP | 2d22h |
| firstpod-service | ClusterIP | 10.111.36.119 | <none>      | 80/TCP       | 3d20h |
| httpd-service    | NodePort  | 10.99.142.63  | <none>      | 80:30081/TCP | 2m22s |

→now Create the Pod in the test Namespace:

--first create test name space:

```
root@master:~# kubectl create ns test  
namespace/test created
```

--now create a pod in test ns

```
root@master: ~  
apiVersion: v1  
kind: Pod  
metadata:  
  name: curl-pod  
  namespace: test  
spec:  
  containers:  
  - name: curl  
    image: curlimages/curl  
    command:  
    - sleep  
    - "3600"
```

--run the yaml

```
root@master:~# kubectl apply -f curl-pod.yaml  
pod/curl-pod created
```

→now check the connectivity:

--in local

```
root@master:~# curl 3.148.252.247:30081  
<html><body><h1>It works!</h1></body></html>
```

-- we can exec into the curl-pod and test the communication with httpdpod by the following command:

```
kubectl exec -it curl-pod -n test -- /bin/sh
```

now check with curl DNS

```
curl httpd-service.default.svc.cluster.local
```

```
~ $ curl httpd-service.default.svc.cluster.local  
<html><body><h1>It works!</h1></body></html>
```

5)Apply a LimitRange with a max limit/request ratio of 2 for memory in the performance-environment namespace, and test by creating a pod with mismatched resource requests and limits.

→create a namespace:

```
root@master:~# kubectl create namespace performance-environment
namespace/performance-environment created
```

Check ns

```
root@master:~# kubectl get ns
NAME                STATUS              AGE
default             Active              4d1h
dev-environment     Terminating       138m
kube-node-lease     Active              4d1h
kube-public         Active              4d1h
kube-system         Active              4d1h
performance-environment Active              17s
```

→create yaml

```
root@master: ~
apiVersion: v1
kind: LimitRange
metadata:
  name: memory-limit-range
  namespace: performance-environment
spec:
  limits:
    - max:
        memory: "4Gi" # Maximum memory limit for any pod in this namespace
      min:
        memory: "512Mi" # Minimum memory limit for any pod in this namespace
      default:
        memory: "1Gi" # Default memory request if not specified
      defaultRequest:
        memory: "1Gi" # Default memory request for the pod
      type: Container
    - maxLimitRequestRatio:
        memory: "2" # Max memory limit to request ratio (limit can't be more than 2x the request)
      type: Pod
```

--run the yaml

```
root@master:~# kubectl apply -f limitrange-memory.yaml
limitrange/memory-limit-range created
```

--check describe

**kubectl describe ns performance-environment**

```
root@master:~# kubectl describe ns performance-environment
Name:          performance-environment
Labels:        kubernetes.io/metadata.name=performance-environment
Annotations:   <none>
Status:        Active
```

No resource quota.

| Resource Limits |          |       |     |                 |               |                         |
|-----------------|----------|-------|-----|-----------------|---------------|-------------------------|
| Type            | Resource | Min   | Max | Default Request | Default Limit | Max Limit/Request Ratio |
| Container       | memory   | 512Mi | 4Gi | 1Gi             | 1Gi           | -                       |
| Pod             | memory   | -     | -   | -               | -             | 2                       |



## → testing by Creating a Pod with Mismatched Resource Requests and Limits

Yaml file:

```
root@master: ~  
apiVersion: v1  
kind: Pod  
metadata:  
  name: test-pod  
  namespace: performance-environment  
spec:  
  containers:  
  - name: test-container  
    image: nginx  
    resources:  
      requests:  
        memory: "1Gi" # Memory request  
      limits:  
        memory: "3Gi" # Memory limit (3Gi is greater than 2x the memory request, so it will be rejected)
```

--run the yaml

**kubectl apply -f test-pod.yaml**

--not created pod got an error because the pod was rejected due to the LimitRange enforcement

```
root@master:~# kubectl apply -f test-pod.yaml  
Error from server (Forbidden): error when creating "test-pod.yaml": pods "test-pod" is forbidden: memory max limit to request ratio per Pod is 2, but provided ratio is 3.000000
```

```
root@master:~# kubectl apply -f test-pod.yaml  
Error from server (Forbidden): error when creating "test-pod.yaml": pods "test-pod" is forbidden: memory max limit to request ratio per Pod is 2, but provided ratio is 3.000000
```