

K8s 08

1) Create and Test a Kubernetes Pod with an EmptyDir Volume

--Yaml file to create a Kubernetes Pod with an EmptyDir Volume.

Here We create volume inside a pod to store data related to container.

```
root@master: ~  
apiVersion: v1  
kind: Pod  
metadata:  
  name: firstpod  
spec:  
  containers:  
  - name: firstcontainer  
    image: nginx  
    volumeMounts:  
    - mountPath: /data #Directory inside container  
      name: first-volume #any logical name  
  volumes:  
  - name: first-volume  
    emptyDir: {} #Blank Object
```

--run the yaml then pod got created

```
root@master:~# vi emptydir.yaml  
root@master:~# kubectl create -f emptydir.yaml  
pod/firstpod created  
root@master:~# kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
firstpod      1/1     Running   0           7s
```

Now test:

--now login to the firstpod

Then do ls now can see data directory

```
root@master:~# kubectl exec -it firstpod bash  
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.  
root@firstpod:~# ls  
bin boot data dev docker-entrypoint.d docker-entrypoint.sh etc home lib lib64 media mnt opt proc root run sbin srv sys tmp usr var
```

Now cd data and create some files in that directory

```
root@firstpod:~# cd data/  
root@firstpod:/data# touch a b c d e f
```

Stop nginx service to terminate your container

```
root@firstpod:/data# cd  
root@firstpod:~# service nginx stop  
command terminated with exit code 137
```

--now exit and check pods again then we can see pod will create a new container

and try to login into container check status

```
root@master:~# exit  
logout  
ubuntu@master:~$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
firstpod      1/1     Running   1 (106s ago)  10m  
ubuntu@master:~$
```

Again login and check also Files in directory /data will be visible

Note:

Problem with empty dir is if the pod got deleted then we will be losing all our data.

2) Configure a HostPath Volume in Kubernetes and Validate Data Persistence

--Yaml file to create a Kubernetes Pod with an HostPath Volume.

Here We create volume on hostpath, means volume will be created outside pod.

```
root@master: ~
apiVersion: v1
kind: Pod
metadata:
  name: firstpod
spec:
  containers:
  - name: firstcontainer
    image: nginx
    volumeMounts:
    - mountPath: /data #Directory inside container
      name: first-volume #any logical name
  volumes:
  - name: first-volume
    hostPath:
      path: /tmp/data # Path inside host machine (Minikube)
```

-run yaml then pod got created

```
root@master:~# vi hostpath.yaml
root@master:~# kubectl create -f hostpath.yaml
pod/firstpod created
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
firstpod      1/1     Running   0           8s
```

-then check on which node pod got created

```
root@master:~# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE              NOMINATED NODE   READINESS GATES
firstpod      1/1     Running   0           37s   192.168.201.231  ip-172-31-13-158  <none>           <none>
```

Now test:

-Login to container and create some random files in /data location.

```
root@master:~# kubectl exec -it firstpod bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@firstpod:/# ls
bin boot data dev docker-entrypoint.d docker-entrypoint.sh etc home lib lib64 media mn
root@firstpod:/# cd data/
root@firstpod:/data# touch created_inside_pod
root@firstpod:/data# ls
created_inside_pod
root@firstpod:/data#
```

Check in worker-02

```
root@worker-02:~# ls
created_inside_pod  snap
```

Now delete the pod and create a new pod.

```
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
firstpod      1/1     Running   0           19m
root@master:~# kubectl delete pods --all
pod "firstpod" deleted
```

Now create new pod

```
root@master:~# kubectl create -f hostpath.yaml
pod/firstpod created
```

Login to new pod and checked if files are available in data directory.

```
root@master:~# kubectl exec -it firstpod bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@firstpod:/# cd data/
root@firstpod:/data# ls
created_inside_pod
root@firstpod:/data# |
```

Note:

If we have multiple nodes then we will other nodes will not be able to access the volume created on node.

3) Deploy an Amazon EBS Volume Using Persistent Volume and Persistent Volume Claim (PVC)

Here If we have multinode k8s cluster, then in this case we need to keep our volume outside the cluster.

If our pod is created on other node then the volume should also be moved to that node.

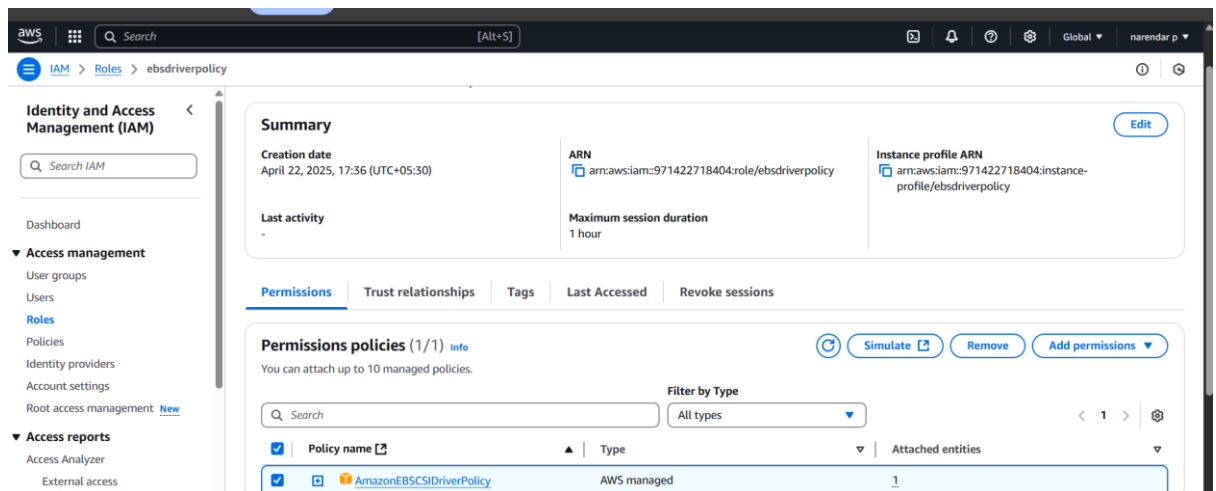
--create an Amazon EBS Volume

The screenshot displays the AWS Management Console interface. On the left, a navigation menu includes sections like Instance Types, Elastic Block Store, and Network & Security. The main content area is titled 'Volumes (1/4)' and shows a table of EBS volumes. The table has columns for Throughput, Snapshot ID, Created, Availability Zone, Volume state, Alarm status, and Attached resources. Below the table, the details for a specific volume (ID: vol-00efe250cc9de3552) are shown, including its size (20 GiB), type (gp3), IOPS (3000), and status (Available).

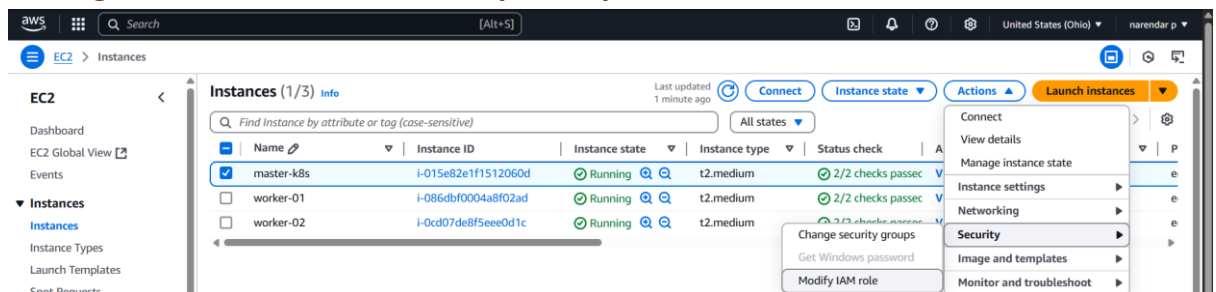
Throughput	Snapshot ID	Created	Availability Zone	Volume state	Alarm status	Attached resources
125	snap-00576a2...	2025/04/13 08:50 GMT+5...	us-east-2a	In-use	No alarms	i-015e82e1f1512060d
125	-	2025/04/22 16:48 GMT+5...	us-east-2a	Available	No alarms	-
125	snap-00576a2...	2025/04/13 08:50 GMT+5...	us-east-2a	In-use	No alarms	i-0cd07de8f5ee0d1c
125	snap-00576a2...	2025/04/13 08:50 GMT+5...	us-east-2a	In-use	No alarms	i-086dbf0004a8f02ad

Volume ID: vol-00efe250cc9de3552	
Details	Status checks
Volume ID vol-00efe250cc9de3552	Status check Okay
AWS Compute Optimizer finding Opt-in to AWS Compute Optimizer for recommendations. Learn more	Throughput 125
Fast snapshot restored No	Multi-Attach enabled No
Size 20 GiB	
Volume state Available	
Availability Zone us-east-2a	
Type gp3	
IOPS 3000	
Created Tue Apr 22 2025 16:48:45 GMT+0530	


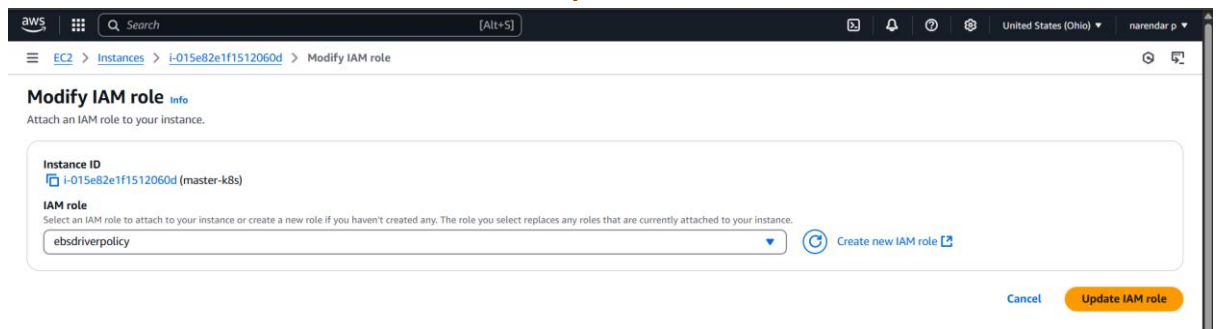
--create an iam role



--now go to ec2 select master-security-modifyIAM role



Then select created iam role and click on update iam role



Do this worker nodes also

--now install EBS CSI driver on k8s masternode by using below url

git clone <https://github.com/kubernetes-sigs/aws-ebs-csi-driver.git>

```
root@master:~# git clone https://github.com/kubernetes-sigs/aws-ebs-csi-driver.git
Cloning into 'aws-ebs-csi-driver'...
remote: Enumerating objects: 33956, done.
remote: Counting objects: 100% (4718/4718), done.
remote: Compressing objects: 100% (944/944), done.
remote: Total 33956 (delta 4010), reused 3778 (delta 3772), pack-reused 29238 (from 2)
Receiving objects: 100% (33956/33956), 28.60 MiB | 2.98 MiB/s, done.
Resolving deltas: 100% (19655/19655), done.
```

Then execute these

cd aws-ebs-csi-driver

kubectl apply -k deploy/kubernetes/overlays/stable/

kubectl get pods -n kube-system #To verify the driver is installed and running.

we can see ebs pods are created

```
root@master:~# cd aws-ebs-csi-driver
root@master:~/aws-ebs-csi-driver# kubectl apply -k deploy/kubernetes/overlays/stable/
serviceaccount/ebs-csi-controller-sa created
serviceaccount/ebs-csi-node-sa created
role.rbac.authorization.k8s.io/ebs-csi-leases-role created
clusterrole.rbac.authorization.k8s.io/ebs-csi-node-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-attacher-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-provisioner-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-resizer-role created
clusterrole.rbac.authorization.k8s.io/ebs-external-snapshotter-role created
rolebinding.rbac.authorization.k8s.io/ebs-csi-leases-rolebinding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-attacher-binding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-node-getter-binding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-provisioner-binding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-resizer-binding created
clusterrolebinding.rbac.authorization.k8s.io/ebs-csi-snapshotter-binding created
deployment.apps/ebs-csi-controller created
poddisruptionbudget.policy/ebs-csi-controller created
daemonset.apps/ebs-csi-node created
csidriver.storage.k8s.io/ebs.csi.aws.com created
root@master:~/aws-ebs-csi-driver# kubectl get pods -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
calico-kube-controllers-658d97c59c-xxs7b  1/1     Running   8 (3h3m ago)  8d
calico-node-496nd                      1/1     Running   0           8d
calico-node-cp9gc                      1/1     Running   0           8d
calico-node-xsfcf                      1/1     Running   8 (3h3m ago)  8d
coredns-76f75df574-6lts2              1/1     Running   7 (3h3m ago)  7d3h
coredns-76f75df574-nvn9b              1/1     Running   0           177m
ebs-csi-controller-6d89bf7f66-12ck6     6/6     Running   0           64s
ebs-csi-controller-6d89bf7f66-q5n26     6/6     Running   0           64s
ebs-csi-node-2h7hj                     3/3     Running   0           64s
ebs-csi-node-cxnvs                     3/3     Running   0           64s
ebs-csi-node-wvcf9                     3/3     Running   0           64s
etcd-master                            1/1     Running   1 (159m ago)  8d
kube-apiserver-master                  1/1     Running   1 (159m ago)  8d
kube-controller-manager-master         1/1     Running   0           8d
kube-proxy-lcq4p                       1/1     Running   8 (3h3m ago)  8d
kube-proxy-m682x                      1/1     Running   0           8d
kube-proxy-pgkkc                       1/1     Running   8 (3h3m ago)  8d
kube-scheduler-master                  1/1     Running   0           8d
metrics-server-75bf97fcc9-qmn68       0/1     Running   1 (3h3m ago)  25h
root@master:~/aws-ebs-csi-driver#
```

--Now we need to create one Persistent volume for our EBS

```
root@master: ~
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-ebs-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  storageClassName: aws-ebs
  awsElasticBlockStore:
    volumeID: vol-00efe250cc9de3552      #EBS volume ID
    fsType: ext4
```

Run yaml then pv created

```
root@master:~# kubectl create -f pv-ebs.yaml
persistentvolume/my-ebs-pv created
root@master:~#
```

```
root@master:~# kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM   STORAGECLASS   VOLUMEATTRIBUTESCLASS   REASON   AGE
my-ebs-pv  10Gi       RWO            Retain           Available           aws-ebs         <unset>                 64s
```

--now create PVC

```
root@master: ~
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-ebs-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  storageClassName: aws-ebs
```

Now run yaml then pvc got created

```
root@master:~# vi pvc-ebs.yaml
root@master:~# kubectl create -f pvc-ebs.yaml
persistentvolumeclaim/my-ebs-pvc created
```

```
root@master:~# kubectl get pvc
NAME          STATUS    VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   VOLUMEATTRIBUTESCLASS   AGE
my-ebs-pvc    Bound     my-ebs-pv 10Gi       RWO           aws-ebs        <unset>                 61s
```

--Now Create pod with PVC

```
root@master: ~
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-container
      image: nginx
      volumeMounts:
        - name: data-volume
          mountPath: /data          #data where we can write data
  volumes:
    - name: data-volume
      persistentVolumeClaim:
        claimName: my-ebs-pvc      #PVC Name
```

Run the yaml then pod got created

```
root@master:~# kubectl apply -f pod.yaml
pod/my-pod created
```

```
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
firstpod      1/1     Running   0           93m
my-pod        1/1     Running   0           46s
```

Now test:

Check pod is created on which node

```
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                NOMINATED NODE   READINESS GATES
firstpod      1/1     Running   0           96m   192.168.201.232  ip-172-31-13-158   <none>           <none>
my-pod        1/1     Running   0           3m44s 192.168.201.238  ip-172-31-13-158   <none>           <none>
```

Now login to pod

Cd data

Crete one file and exit

```
root@master:~# kubectl exec -it my-pod bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@my-pod:~# cd data/
root@my-pod:/data# ls
lost+found
root@my-pod:/data# touch create_in_worker-02-pod
root@my-pod:/data# exit
```

Once pod is created then we can check the volume as it will show in use status.

Delete the pod and try to see the status of volume

```
root@master:~# kubectl delete pods --all
pod "firstpod" deleted
```

```
root@master:~# kubectl get pvc my-ebs-pvc
NAME          STATUS  VOLUME  CAPACITY  ACCESS  MODES  STORAGECLASS  VOLUMEATTRIBUTESCLASS  AGE
my-ebs-pvc    Bound   my-ebs-pv  10Gi      RWX     RWO      aws-ebs        <unset>                 17m
```

Create new pod again and this time if it is scheduled on another node(worker-02) then volume will be attached to that node.

```
root@master:~# kubectl apply -f pod.yaml
pod/my-pod created
```

```
root@master:~# kubectl get pods -o wide
NAME    READY  STATUS   RESTARTS  AGE  IP            NODE          NOMINATED NODE  READINESS GATES
my-pod  1/1    Running  0          38s  192.168.201.237  ip-172-31-13-158  <none>          <none>
```

Now login and check volume attached to that node.

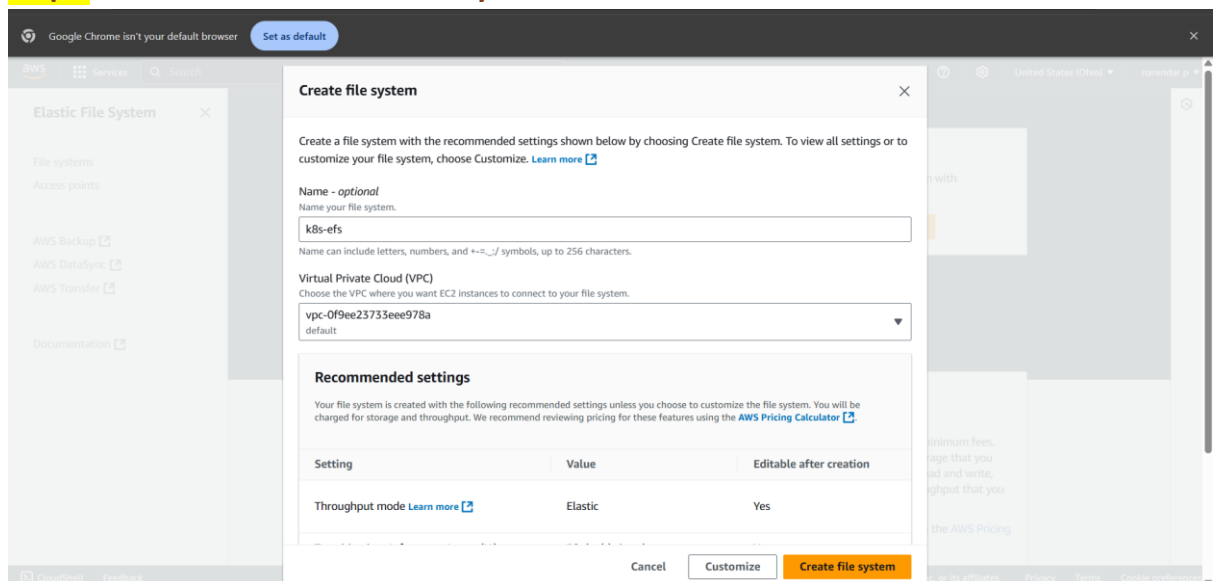
```
root@master:~# kubectl exec -it my-pod bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@my-pod:~# cd data/
root@my-pod:/data# ls
create_in_worker-02-pod  lost+found
root@my-pod:/data#
```

Note:

Here we can map multiple pods to the volumes but we cannot map two pods to same volume at same point of time.

4) Set Up an Amazon EFS Volume and Attach it to Multiple Pods

Step 1: Create an Amazon EFS File System



Success!
File system (fs-0e477bc908d5f2534) is available.

[View file system](#)

Amazon EFS > File systems

File systems (1)

Filter by property values

	Name	File system ID	Encryption	Total size	Size in Standard	Size in IA	Size in Archive	Provisioned Throughput (MiB/s)
<input type="radio"/>	k8s-efs	fs-0e477bc908d5f2534	Encrypted	6.00 KiB	6.00 KiB	0 Bytes	0 Bytes	-

Amazon EFS > File systems > fs-0e477bc908d5f2534

k8s-efs (fs-0e477bc908d5f2534)

[Delete](#) [Attach](#)

[Edit](#)

General

Amazon resource name (ARN)
arn:aws:elasticfilesystem:us-east-2:971422718404:file-system/fs-0e477bc908d5f2534

Performance mode
General Purpose

Throughput mode
Elastic

Lifecycle management
Transition into Infrequent Access (IA): 30 day(s) since last access
Transition into Archive: 90 day(s) since last access
Transition into Standard: None

Availability zone
Regional

Automatic backups
Enabled

Encrypted
fe5845d1-d7f2-4e95-af59-b09702021fa4 (aws/elasticfilesystem)

File system state
Available

DNS name
fs-0e477bc908d5f2534.efs.us-east-2.amazonaws.com

Replication overwrite protection
Enabled

Step 2: Create an Access Point

Amazon EFS > Access points > Create

Create access point for fs-0e477bc908d5f2534

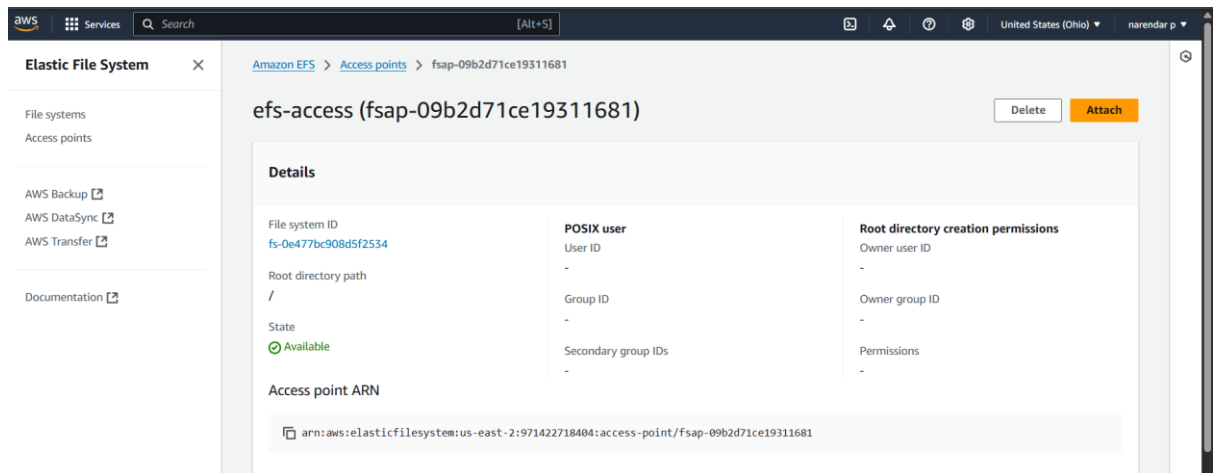
An access point is an application-specific entry point into an EFS file system that makes it easier to manage application access to shared datasets. [Learn more](#)

Details

File system
Choose the file system to which your access point is associated.
fs-0e477bc908d5f2534

Name - optional
efs-access
Name can include letters, numbers, and +,=,_,/ symbols, up to 256 characters.

Root directory path - optional
Connections use the specified path as the file system's virtual root directory [Learn more](#)
/
Example: "/foo/bar"



Step 3: Install the Amazon EFS CSI Driver

```
kubectl apply -k "github.com/kubernetes-sigs/aws-efs-csi-driver/deploy/kubernetes/overlays/stable/ecr/?ref=release-1.7"
```

```
root@master:~# kubectl apply -k "github.com/kubernetes-sigs/aws-efs-csi-driver/deploy/kubernetes/overlays/stable/ecr/?ref=release-1.7"
# Warning: 'bases' is deprecated. Please use 'resources' instead. Run 'kustomize edit fix' to update your Kustomization automatically.
serviceaccount/efs-csi-controller-sa created
serviceaccount/efs-csi-node-sa created
clusterrole.rbac.authorization.k8s.io/efs-csi-external-provisioner-role created
clusterrole.rbac.authorization.k8s.io/efs-csi-external-provisioner-role-describe-secrets created
clusterrole.rbac.authorization.k8s.io/efs-csi-node-role created
rolebinding.rbac.authorization.k8s.io/efs-csi-provisioner-binding-describe-secrets created
clusterrolebinding.rbac.authorization.k8s.io/efs-csi-node-binding created
clusterrolebinding.rbac.authorization.k8s.io/efs-csi-provisioner-binding created
deployment.apps/efs-csi-controller created
daemonset.apps/efs-csi-node created
csidriver.storage.k8s.io/efs.csi.aws.com created
```

Step 4: Create a StorageClass

Yaml file

```
root@master: ~
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: efs-sc
provisioner: efs.csi.aws.com
```

Run yaml

```
root@master:~# kubectl apply -f efs-sc.yaml
storageclass.storage.k8s.io/efs-sc created
```

Step 5: Create a Persistent Volume (PV)

Yaml file

```
root@master: ~
apiVersion: v1
kind: PersistentVolume
metadata:
  name: efs-pv
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  storageClassName: efs-sc
  csi:
    driver: efs.csi.aws.com
    volumeHandle: fs-0e477bc908d5f2534::fsap-09b2d71ce19311681 # Format: <FileSystemId>::<AccessPointId>
```

Run yaml then PV got created

```
root@master:~# kubectl apply -f efs-pv.yaml
persistentvolume/efs-pv created
root@master:~# kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	VOLUMEATTRIBUTESCLASS	REASON	AGE
efs-pv	5Gi	RWX	Retain	Available		efs-sc	<unset>		19s

Step 6: Create a Persistent Volume Claim (PVC)

Yaml file

```
root@master: ~
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: efs-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: efs-sc
  resources:
    requests:
      storage: 5Gi
```

Run the yaml file then pvc got created

```
root@master:~# kubectl apply -f efs-pvc.yaml
persistentvolumeclaim/efs-pvc created
```

```
root@master:~# kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	VOLUMEATTRIBUTESCLASS	AGE
efs-pvc	Bound	efs-pv	5Gi	RWX	efs-sc	<unset>	46s

Step 7: Use PVC in Multiple Pods

```
root@master: ~
apiVersion: v1
kind: Pod
metadata:
  name: pod1
spec:
  containers:
    - name: app
      image: busybox
      command: [ "sleep", "3600" ]
      volumeMounts:
        - name: efs-vol
          mountPath: /mnt/efs
  volumes:
    - name: efs-vol
      persistentVolumeClaim:
        claimName: efs-pvc
```

Run yaml file

```
kubectl apply -f pod.yaml
root@master:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
efs-pod	1/1	Running	0	47m

Test now:

Now login to pod and create some random files in efs-mount directory.

```
root@master:~# kubectl exec -it efs-pod -- bash
root@efs-pod:/# cd /efs-mount
```

```
root@master:~# kubectl exec -it efs-pod -- bash
root@efs-pod:/# cd /efs-mount
echo "Hello from EFS!" > hello.txt
touch random-file-{1..5}.txt

root@efs-pod:/efs-mount# ls
hello.txt random-file-1.txt random-file-2.txt random-file-3.txt random-file-4.txt random-file-5.txt
```

Now Create one more pod and check if the files are available or not

In above yaml file just edit name to pod2 and apply and create one more pod

```
root@master:~# kubectl apply -f pod.yaml
pod/efs-pod2 created
```

Again login and check the files

```
root@master:~# kubectl exec -it efs-pod2 -- bash
root@efs-pod2:/# ls
bin boot dev docker-entrypoint.d docker-entrypoint.sh efs-mount etc home lib lib64 media mnt opt proc root run
root@efs-pod2:/# cd efs-mount/
root@efs-pod2:/efs-mount# ls
hello.txt random-file-1.txt random-file-2.txt random-file-3.txt random-file-4.txt random-file-5.txt
root@efs-pod2:/efs-mount#
```

We can now connect multiple pods to same volume

Note:

Here we can map multiple pods to same EFS volume.

EFS cost is 3time higher to EBS.

5) Implement and Test Liveness and Readiness Probes in a Kubernetes

Pod

---yaml file to create pod with liveness probe

```
root@master: ~
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
  - name: my-app
    image: nginx          #Image Name
    ports:
    - containerPort: 80
    livenessProbe:
      httpGet:
        path: /           #get request or we can use tcpSocket, or exec
        port: 80          #health check path
      initialDelaySeconds: 15 #It will wait 15 second before applying first liveness
      periodSeconds: 10      #Every 10 seconds the liveness will be checked.
```

--run the yaml

```
kubectl apply -f pod.yaml
```

```
root@master:~# vi pod.yaml
root@master:~# kubectl apply -f pod.yaml
pod/my-pod created
```

-pod created

```
root@master:~# kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
my-pod    1/1     Running   0           8s
```

-describe the pod

kubectl describe pod my-pod

we can see liveness of pod

```
Status:          Running
IP:             192.168.201.229
IPs:
  IP: 192.168.201.229
Containers:
  my-app:
    Container ID:  containerd://b1281635a727dc8fca880e4ac7838f38771ae6d95a1ccdcfbd753bf431794ed9
    Image:         nginx
    Image ID:      docker.io/library/nginx@sha256:5ed8fcc66f4ed123c1b2560ed708dc148755b6e4cbd8b943fab094f
    Port:         80/TCP
    Host Port:    0/TCP
    State:        Running
      Started:    Tue, 22 Apr 2025 09:45:42 +0000
    Ready:        True
    Restart Count: 0
    Liveness:     http-get http://:80/ delay=15s timeout=1s period=10s #success=1 #failure=3
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-48ccg (ro)

Events:
  Type     Reason      Age   From              Message
  ----     -
  Normal   Scheduled   30s   default-scheduler Successfully assigned default/my-pod to ip-172-31-13-158
  Normal   Pulling     29s   kubelet           Pulling image "nginx"
  Normal   Pulled      29s   kubelet           Successfully pulled image "nginx" in 215ms (215ms including waiting)
  Normal   Created     29s   kubelet           Created container: my-app
  Normal   Started     29s   kubelet           Started container my-app
```

Now Test Liveness probe of a my-pod

So create one yaml with wrong container ports

```
root@master: ~
apiVersion: v1
kind: Pod
metadata:
  name: first-pod
spec:
  containers:
  - name: my-app
    image: nginx #Image Name
    ports:
    - containerPort: 8080
    livenessProbe:
      httpGet: #get request or we can use tcpSocket, or exec
        path: / #health check path
        port: 8080
      initialDelaySeconds: 5 #It will wait 15 second before applying first liveness
      periodSeconds: 5 #Every 10 seconds the liveness will be checked.
```

--run yaml file

First-Pod created

```
root@master:~# vi pod.yaml
root@master:~# kubectl apply -f pod.yaml
pod/first-pod created
root@master:~# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
first-pod     1/1     Running   0          10s
my-pod        1/1     Running   0          12m
```

-describe first-pod

We can see Container my-app failed liveness probe, will be restarted

```
Events:
  Type     Reason      Age   From              Message
  ----     -
  Normal   Scheduled   42s   default-scheduler Successfully assigned default/first-pod to ip-172-31-13-158
  Normal   Pulled      41s   kubelet           Successfully pulled image "nginx" in 228ms (228ms including waiting)
  Normal   Pulled      22s   kubelet           Successfully pulled image "nginx" in 226ms (226ms including waiting)
  Normal   Pulling     2s (x3 over 42s) kubelet           Pulling image "nginx"
  Warning  Unhealthy   2s (x6 over 32s) kubelet           Liveness probe failed: Get "http://192.168.201.230:8080/": dial tcp 192
used
  Normal   Killing     2s (x2 over 22s) kubelet           Container my-app failed liveness probe, will be restarted
  Normal   Created     1s (x3 over 41s) kubelet           Created container: my-app
  Normal   Started     1s (x3 over 41s) kubelet           Started container my-app
  Normal   Pulled      1s   kubelet           Successfully pulled image "nginx" in 439ms (439ms including waiting)
```

-check pods now

It restarted got failed

```
root@master:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
first-pod	0/1	CrashLoopBackOff	6 (57s ago)	5m22s
my-pod	1/1	Running	0	17m

----yaml file to create pod with liveness probe

```
root@master: ~  
apiVersion: v1  
kind: Pod  
metadata:  
  name: my-pod1  
spec:  
  containers:  
  - name: my-app  
    image: nginx  
    ports:  
    - containerPort: 80  
    readinessProbe:  
      httpGet:  
        path: /index.html  
        port: 80  
      initialDelaySeconds: 10  
      periodSeconds: 5
```

--run yaml file

My-Pod1 created

```
root@master:~# vi pod1.yaml  
root@master:~# kubectl apply -f pod1.yaml  
pod/my-pod1 created  
root@master:~# vi pod1.yaml  
root@master:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
first-pod	0/1	CrashLoopBackOff	7 (2m54s ago)	10m
my-pod	1/1	Running	0	22m
my-pod1	1/1	Running	0	81s

-describe my-pod1

We can see readiness of a pod

```
Containers:  
  my-app:  
    Container ID:   containerd://99f65d1c8c4391ff85aa4369a2d437e0350ebb27d7f6a4c93158b03f0d47d152  
    Image:          nginx  
    Image ID:       docker.io/library/nginx@sha256:5ed8fcc66f4ed123c1b2560ed708dc148755b6e4cbd8b943fab094f2c6bfa91e  
    Port:           80/TCP  
    Host Port:      0/TCP  
    State:          Running  
      Started:      Tue, 22 Apr 2025 10:06:50 +0000  
    Ready:          True  
    Restart Count:  0  
    Readiness:      http-get http://:80/index.html delay=10s timeout=1s period=5s #success=1 #failure=3  
    Environment:    <none>  
    Mounts:  
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-955qm (ro)
```

Now test

For that Login to the pod and delete the index.html file

```
root@master:~# kubectl exec -it my-pod1 bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
root@my-pod1:/# cd /usr/share/nginx/html
root@my-pod1:/usr/share/nginx/html# rm -rf *
root@my-pod1:/usr/share/nginx/html#
```

Now check pods can see pod is not ready not serve the traffic

```
root@master:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
first-pod	0/1	CrashLoopBackOff	9 (5m9s ago)	18m
my-pod	1/1	Running	0	30m
my-pod1	0/1	Running	0	9m21s

-describe my-pod1

Readiness probe failed

```
Events:
  Type     Reason      Age           From          Message
  ----     -
Normal    Scheduled   9m54s        default-scheduler   Successfully assigned default/my-pod1 to ip-172-31-13-158
Normal    Pulling     9m53s        kubelet        Pulling image "nginx"
Normal    Pulled      9m53s        kubelet        Successfully pulled image "nginx" in 207ms (207ms including waiting)
Normal    Created     9m53s        kubelet        Created container: my-app
Normal    Started     9m53s        kubelet        Started container: my-app
Warning   Unhealthy   4s (x18 over 84s)  kubelet        Readiness probe failed: HTTP probe failed with statuscode: 404
```